

Lab 08 – Mini-Lecture Microbenchmarks

Prof. Brendan Kochunas

NERS/ENGR 570 - Methods and Practice of Scientific Computing (F20)



COLLEGE OF ENGINEERING

NUCLEAR ENGINEERING & RADIOLOGICAL SCIENCES

UNIVERSITY OF MICHIGAN

Outline

- What are microbenchmarks?
- The Membench Experiment
 - Design
 - Analysis
- Results of benchmark from various machines
- Instructions
- 1-slide on SIMD

Learning Objectives: By the end of Today's Lab you should be able to

- (Knowledge) understand how to profile and analyze a memory structure of a single node/computer.
- (Skill) generate results for membench on Great Lakes.
- (Knowledge) quickly look for and identify relevant benchmarks for measuring system performance parameters
- (Knowledge) describe what a mini-app is.

Microbenchmarks

- Purpose

- Measure the performance of the full stack of
 - Source code
 - Hardware
 - OS
 - Compiler
 - Compiler options
- validate performance models

- Scope

- Quite small, effectively they are single procedures or short snippets of code.


- Examples

- Membench
- [STREAM](#)
- [OSU MPI Micro-Benchmarks](#)
- [OpenMP Benchmarks](#)
- PSNAP
- IOR
- MDTEST
- PYNAMIC
- SMB
- UPC FT
- MPIMEMU
- [NERSC-8 Benchmark Suite](#)



Extension to Mini-Apps

- Purpose
 - Evaluate the performance and scalability of a particular algorithm or numerical method
- Scope
 - Small programs that mimic the computational kernels of every step, but do not use “real data”
- Examples
 - [Exascale Computing Project](#)
 - AMG, Ember, ExaMiniMD, Laghos, MACSio, miniAMR, miniQMC, miniVite, **NEKbone**, **PICSARlite**, SW4lite, SWFFT, Thornado-mini, **XSbench**
 - miniGAN, miniRL, CRADL, Cosmoflow-Benchmark, Mlperf-DeepCAM
 - [NERSC-8 Benchmarks](#)
 - [LANL Proxy Apps](#)

Memory Hierarchy



On chip




	Register	L1	L2	L3	DRAM	Disk	Tape
Size	< 1 KB	~1KB	1 MB	10's MB	1-100's GB	TB	PB
Speed	< 1ns	<1 ns	~1 ns	~1-10 ns	10-100 ns	10 ms	~10s

Cache Basics

- Cache is fast (expensive) memory which keeps copy of data in main memory;
 - Typically it is hidden from software (e.g. no standard way of programming directly)
 - Simplest example: data at memory address xxxxx1101 is stored at cache location 1101
- Cache hit: in-cache memory access—cheap
- Cache miss: non-cached memory access—expensive
 - Need to access next, slower level of cache
- Cache line length: # of bytes loaded together in one entry
 - Ex: If either xxxxx1100 or xxxxx1101 is loaded, both are
- Associativity
 - direct-mapped: only 1 address (line) in a given range in cache
 - Data stored at address xxxxx1101 stored at cache location 1101, in 16 word cache
 - n -way: $n \geq 2$ lines with different addresses can be stored
 - Up to $n \leq 16$ words with addresses xxxxx1101 can be stored at cache location 1101 (so cache can store $16n$ words)

Techniques for exposing details of Memory




- Determine memory access times experimentally using “micro”-benchmarks
 - Membench (Saavedra-Barrera), STREAM, others...


- for array A of length L from 4KB to 8MB by 2x
for stride s from 4 Bytes (1 word) to L/2 by 2x
time the following loop
(repeat many times and average)
for i from 0 to L **by s**
load A[i] from memory (4 Bytes)

1 experiment

Membench: What to Expect

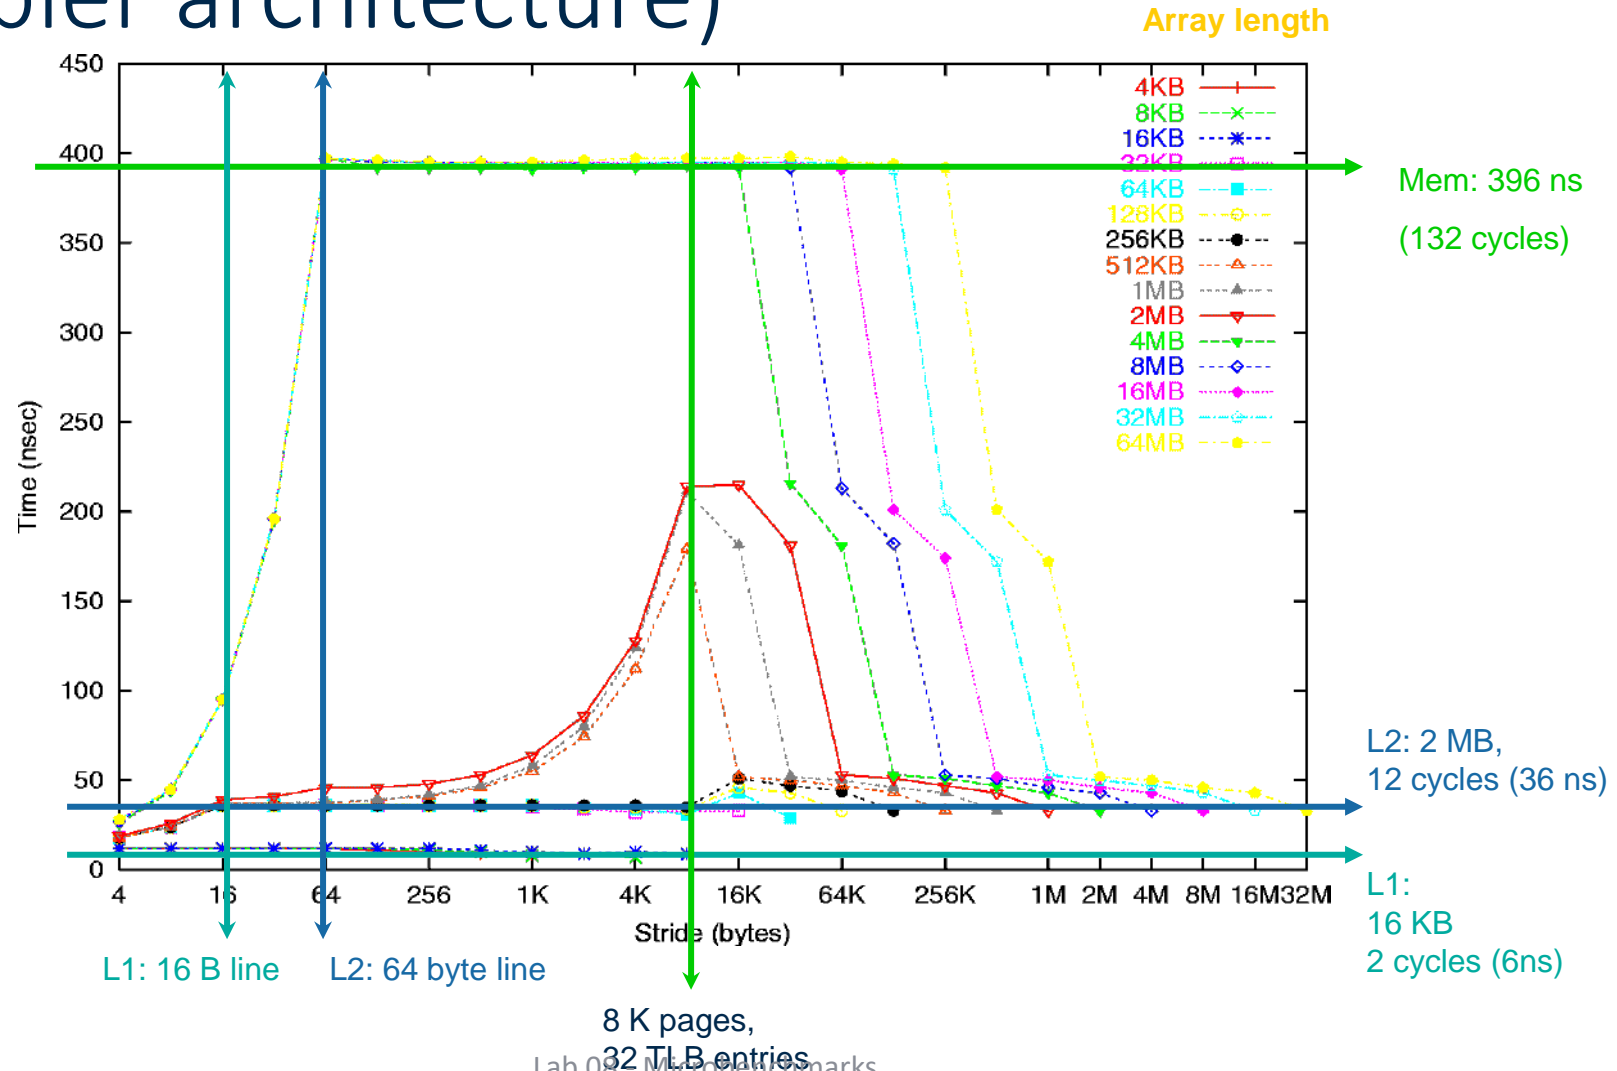


- Consider the average cost per load
 - Plot one line for each array length: time vs. stride
 - If array is smaller than a given cache all accesses will hit (after first run)
 - Small stride is best: e.g. if cache line holds 4 words, expect $\frac{1}{4}$ miss
 - Picture assumes one-level cache
 - More difficult to measure on modern processors due to more complex memory systems



Results for a 25 year old computer (e.g. simpler architecture)

See
www.cs.berkeley.edu/~yelick/arvindk/t3d-isca95.ps for more details



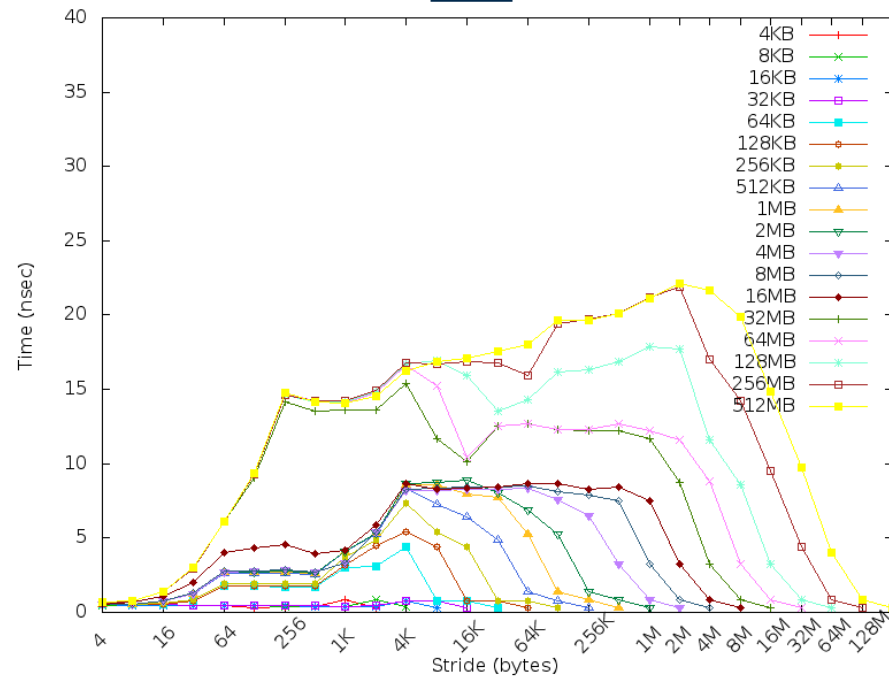
Example Membench Results
for Sun Ultra-2i, 333 MHz

Results from 10-year old Workstation




More Recent Results

Flux



Great Lakes




For the experiment you are expected to:

- Modify the program to consider arrays of up to 1GB.
- Modify the program to fix any output formatting
- Modify the provided data processing scripts to generate the plots for the given number of array trials.
 - Two scripts are provided: a script for gnuplot and one for Python
- Run the experiment on a fully dedicated compute node.

For the Data Analysis

- Identify as many features as you can from the plot you generate and annotate this plot.
 - Similar to slide 8
- At a minimum, you are expected to identify:
 - L1 cache size, line size, and access time.
 - L2 cache size, line size, and access time.
 - L3 cache size, line size, and access time.
 - Main memory access time.
- Provide a reasoned explanation for why your data suggests that the <memory QOI> is <value>.
- Think about how you can verify the hardware specifications. Can you find info about the hardware and see if it matches the features you identify in your plot?
- For data that you are not quite sure how to explain, provide possible explanations as to what could be going on.
- If you are uncertain about values derived from your plot(s), or if the measurements contain some noise, it is ok to provide a range of values.
- You may have to create multiple plots to zoom in on certain features.

Note about node hardware



Single Instruction Multiple Data

- One operation produces multiple results
- Implemented as SSE assembly instructions by compiler
 - SSE = Streaming SIMD Instructions
 - Several standards SSE (128-bit), SSE2, ... SSE4, AVX (256-bit), AVX2 (512-bit)
- Operate on anything that fits into x bytes (e.g. 16 bytes)
 - Operations include add, multiply, etc.
- Challenges
 - Need to be contiguous in memory *and aligned*
 - Some instructions are needed to move data around from one register to another

