



COLLEGE OF ENGINEERING  
NUCLEAR ENGINEERING & RADIOLOGICAL SCIENCES  
UNIVERSITY OF MICHIGAN

# Lecture 7

## Solving Linear Systems and Scientific Computing Libraries

Prof. Brendan Kochunas  
9/25/2019

NERS 590-004



# Outline

- Short Recap of Last Lecture
- Classical Iterative Methods
- Multigrid Methods
- Krylov Methods
- Scientific Computing Libraries



## Learning Objectives

- Understand the differences of various approaches to solving linear systems
- Understand how to determine what method might be best for your application
- Become aware of popular, high quality scientific computing libraries



# Basic Linear Algebra Operations

## Residual and Norms of Vectors

$$\mathbf{r} = \mathbf{Ax} - \mathbf{b} \quad \text{residual}$$

$$\|\mathbf{r}\|_1 = \sum_i |r_i| \quad \text{1-norm}$$

$$\|\mathbf{r}\|_2 = \sqrt{\sum_i r_i^2} \quad \text{2-norm ("average error")}$$

$$\|\mathbf{r}\|_\infty = \max_i (|r_i|) \quad \infty\text{-norm ("max local error")}$$

$$\|\mathbf{r}\|_p = \left( \sum_i |r_i|^p \right)^{1/p} \quad p\text{-norm}$$

## Inner/Dot Product (vector-vector multiply)

$$\mathbf{u}^T \cdot \mathbf{v} = \sum_i u_i v_i$$

## Matrix-vector Multiply

$$\mathbf{Ax} = \mathbf{b} \rightarrow b_i = \sum_j a_{i,j} x_j$$

## Matrix-Matrix Multiply

$$\mathbf{AB} = \mathbf{C} \rightarrow c_{i,j} = \sum_k a_{i,k} b_{k,j}$$



# LU and QR Decompositions

## LU Decomposition

$$\mathbf{A} = \mathbf{L}\mathbf{U}$$

Useful for solving linear systems

$$\mathbf{A}\mathbf{x} = \mathbf{b}$$

$$\mathbf{L}\mathbf{U}\mathbf{x} = \mathbf{b} \quad \mathbf{L}^{-1}\mathbf{L}\mathbf{U}\mathbf{x} = \mathbf{L}^{-1}\mathbf{b}$$

$$\mathbf{L}\mathbf{y} = \mathbf{b} \quad \text{Forward elimination}$$

$$y_i = \frac{b_i - \sum_{j=1, i} \ell_{i,j} y_j}{\ell_{i,i}}$$

$$\mathbf{U}\mathbf{x} = \mathbf{y} \quad \text{Backward Substitution}$$

$$x_i = \frac{y_i - \sum_{j=n, i, -1} u_{i,j} x_j}{u_{i,i}}$$

## QR Decomposition

$$\mathbf{A} = \mathbf{Q}\mathbf{R}$$

Useful for obtaining orthonormal basis

Obtaining by Gram-Schmidt

$$\text{proj}_{\mathbf{a}_{i-1}} \mathbf{a}_i = \frac{\mathbf{a}_{i-1}^T \cdot \mathbf{a}_i}{\mathbf{a}_{i-1}^T \cdot \mathbf{a}_{i-1}} \mathbf{a}_{i-1} \quad \text{project}$$

$$\mathbf{u}_i = \mathbf{a}_i - \text{proj}_{\mathbf{a}_{i-1}} \mathbf{a}_i \quad \text{orthogonalize}$$

$$\mathbf{q}_i = \frac{\mathbf{u}_i}{\|\mathbf{u}_i\|} \quad \text{normalize}$$

$$r_{i,j} = \mathbf{q}_i^T \cdot \mathbf{a}_j$$

# Solving Linear Systems

$$Ax = b$$


→ MATLAB  $x = b \backslash A$   
 ↑  
 "Solve"

## Iterative Methods

- "Classical"
  - Gauss-Seidel / Jacobi
  - Successive Over-Relaxation
- Multigrid (1970's - 1980's)
- Krylov Methods
  - GMRES (1990's)
  - Lanczos & Arnoldi's (1950's)
  - Conjugate Gradient

## DIRECT METHODS

- Gaussian Elim.
- LU decomp.



Direct -  $O(n^3)$

Iterative -  $O(n^2) \times m$


} Symmetric  
 } Non-Symmetric



# Classical Iterative Methods

# Classical Iteration Schemes

$$\mathbf{A} = (\mathbf{L} + \mathbf{D} + \mathbf{U}) =$$



## Jacobi

$$x_i^{(\ell+1)} = -\frac{1}{a_{ii}} \left( \sum_{j=1}^{i-1} l_{ij} x_j^{(\ell)} + \sum_{j=i+1}^n u_{ij} x_j^{(\ell)} \right) x_i^{(\ell)} + b_i$$

## Gauss-Siedel

$$x_i^{(\ell+1)} = -\frac{1}{a_{ii}} \left( \sum_{j=1}^{i-1} l_{ij} x_j^{(\ell+1)} + \sum_{j=i+1}^n u_{ij} x_j^{(\ell)} \right) x_i^{(\ell)} + b_i$$

$$\mathbf{x}^{(\ell+1)} = -\mathbf{D}^{-1}(\mathbf{L} + \mathbf{U})\mathbf{x}^{(\ell)} + \mathbf{D}^{-1}\mathbf{b}$$

$$\mathbf{F} = -\mathbf{D}^{-1}(\mathbf{L} + \mathbf{U})$$

$$\mathbf{c} = \mathbf{D}^{-1}\mathbf{b}$$

$$\mathbf{x}^{(\ell+1)} = -(\mathbf{D} + \mathbf{L})^{-1}\mathbf{U}\mathbf{x}^{(\ell)} + (\mathbf{D} + \mathbf{L})^{-1}\mathbf{b}$$

$$\mathbf{F} = -(\mathbf{D}^{-1} + \mathbf{L})\mathbf{U}$$

$$\mathbf{c} = (\mathbf{D}^{-1} + \mathbf{L})\mathbf{b}$$

$$\boxed{\mathbf{x}^{(\ell+1)} = \mathbf{F}\mathbf{x}^{(\ell)} + \mathbf{c}}$$





## Do they converge?

- Fixed point iteration

$$\mathbf{x}^{(\ell+1)} = \mathbf{F}\mathbf{x}^{(\ell)} + \mathbf{c}$$

- Express iterate as combination of exact solution and error

$$\mathbf{x} + \boldsymbol{\varepsilon}^{(\ell+1)} = \mathbf{F}(\mathbf{x} + \boldsymbol{\varepsilon}^{(\ell)}) + \mathbf{c}$$

- If the method converges then:

$$\lim_{\ell \rightarrow \infty} \boldsymbol{\varepsilon}^{(\ell)} = 0$$



## Condition for Convergence

Evolution of error?

$$\varepsilon^{(l+1)} = F \varepsilon^{(l)}$$

$$\lim_{l \rightarrow \infty} \varepsilon^{(l)} = 0$$

Recall Eigen decomposition  $A = Q \Lambda Q^{-1}$

$$Q Q^{-1} = I \quad A^2 = (Q \Lambda Q^{-1})(Q \Lambda Q^{-1}) = Q \Lambda^2 Q^{-1}$$

$$\varepsilon^{(1)} = F \varepsilon^{(0)}$$

$$\varepsilon^{(2)} = F \varepsilon^{(1)} = F^2 \varepsilon^{(0)}$$

$$\varepsilon^{(l)} = Q \Lambda^l Q^{-1} \varepsilon^{(0)}$$

$$|\lambda_1| > |\lambda_2| > \dots$$



## More about the spectral radius

$|\lambda_1| = |\lambda_{\max}| = \text{spectral radius}$   
 $\rho(F)$

- The spectral radius determines the rate of convergence (for fixed point iteration schemes)

$$\rho(F) < 1 \quad \|Ax\| \leq \|A\| \|x\| \quad \|e^{(k+1)}\| \leq \|F\| \|e^{(k)}\|$$

$$\frac{\|e^{(k+1)}\|}{\|e^{(k)}\|} \leq \|F\| \quad \frac{\|e^{(k+1)}\|}{\|e^{(k)}\|} \approx \rho(F) = \lim_{k \rightarrow \infty} \frac{\|e^{(k+1)}\|}{\|e^{(k)}\|}$$

$$e^{(k)} = r^{(k)} \|Ax^{(k)} - b\| \quad 0.99^2 = 10^{-5} \rightarrow 1,100 \quad 0.3^2 = 10^{-5} \rightarrow \sim 10$$



# Summary of Classical Iteration Schemes

- Implementations are very simple
- Error properties are very well understood
- Generally slowly converging in practical problems
- Good for simple problems



# Multigrid Methods

# Multigrid Methods

- Logical extension to classical methods that arises from error analysis.
  - Consider “shape” of error  
→ frequency transform






Figure 2.2: The modes  $v_j = \sin\left(\frac{jk\pi}{n}\right)$ ,  $0 \leq j \leq n$ , with wavenumbers  $k = 1, 3, 6$ .  
The  $k$ th mode consists of  $\frac{k}{2}$  full sine waves on the interval.

Low Frequency



Classical methods  
very good at  
“smoothing”  
high-frequency errors



Real Problem  
(multiple modes)





Figure 2.9: Weighted Jacobi method with  $\omega = \frac{2}{3}$  applied to the one-dimensional model problem with  $n = 64$  points and with an initial guess consisting of (a)  $\mathbf{w}_3$ , (b)  $\mathbf{w}_{16}$ , and (c)  $(\mathbf{w}_2 + \mathbf{w}_{16})/2$ . The figures show the approximation after one iteration (left side) and after 10 iterations (right side).

## Multigrid Methods (2)

- Central idea of multigrid is to “map” errors onto coarser grids
  - A low-frequency error on a fine-grid is a high-frequency error on a coarse-grid!
- Recipe for Multigrid includes
  - How to map error from fine-grid to coarse-grid?
    - restriction operator (e.g. bi-linear average)
  - How to smooth error on each grid?
    - classical iteration scheme
  - How to correct error in fine-grid from coarse grid?
    - interpolation operator (e.g. linear interpolate)
  - How to traverse grids?






# Multigrid: Restriction and Interpolation


## Restriction

$$I_h^{2h} \mathbf{v}^h = \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 & & & & \\ & 1 & 2 & 1 & & & \\ & & 1 & 2 & 1 & & \\ & & & 1 & 2 & 1 & \\ & & & & 1 & 2 & 1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \\ v_5 \\ v_6 \\ v_7 \end{bmatrix}_h = \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix}_{2h} = \mathbf{v}^{2h}$$



## Interpolation


$$I_{2h}^h \mathbf{v}^{2h} = \frac{1}{2} \begin{bmatrix} 1 & & & & \\ 2 & & & & \\ 1 & 1 & & & \\ & 2 & & & \\ & 1 & 1 & 2 & \\ & & 1 & & 1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix}_{2h} = \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \\ v_5 \\ v_6 \\ v_7 \end{bmatrix}_h = \mathbf{v}^h$$






# Multigrid: Traversing the Grids


## V-cycle



## W-cycle



## Full Multigrid





## Summary of Multigrid

- Very good for elliptic problems
- A type of fixed point iteration
  - May be analyzed via Fourier/Von Neumann Analysis for asymptotic convergence
- Builds on traditional classical fixed point iterative techniques
  - Uses same elements and adds a few more (interpolation/prolongation)
- Lots of parameters in the iteration that can be “tuned”
- Good for structured grids and finite differenced or finite volume disc (e.g. discretized operator is a stencil)
- Can be generalized to algebraic multi-grid (AMG)



# Krylov Methods

Reference: Yousef Saad, *Iterative Methods for Sparse Linear Systems*, SIAM



## Krylov (subspace) Methods

- Use Krylov subspaces

$$K_p(\mathbf{A}, \mathbf{b}) = \text{span}\{\mathbf{b}, \mathbf{A}\mathbf{b}, \mathbf{A}^2\mathbf{b}, \dots, \mathbf{A}^p\mathbf{b}\}$$

- Krylov subspaces are easy and efficient to construct
  - Use only matrix-vector operations (this is good for large sparse systems!)
- Krylov methods also rely on projection
  - Specifically projection to the Krylov subspace

## General Idea: Projection Methods

- Two types: orthogonal and oblique
- For a linear system:  $\mathbf{Ax} = \mathbf{b}$ 
  - $\mathbf{A}$  exists in  $\mathbf{R}^n$  and  $\mathbf{K}$  and  $\mathbf{L}$  are two subspaces within  $\mathbf{R}^n$
  - An approximation of the solution is:  $\tilde{\mathbf{x}} = \mathbf{x}_0 + \delta$
  - which has the residual:  $\mathbf{r} = \mathbf{Ax} - \mathbf{b}$
- Solution by orthogonal projection means

Find  $\tilde{\mathbf{x}} \in \mathbf{x}_0 + \mathbf{K}$  such that  $\mathbf{b} - \mathbf{Ax}^{(\ell)} \perp \mathbf{L}$

$$\mathbf{r}_0 = \mathbf{b} - \mathbf{Ax}_0$$

$$\mathbf{r}_0 = \mathbf{b} - \mathbf{A}(\mathbf{x}_0 + \delta) \perp \mathbf{L}$$

$$\mathbf{r}_0 = \mathbf{b} - \mathbf{A}\delta \perp \mathbf{L}$$

$$(\mathbf{r}_0 - \mathbf{A}\delta)^T \cdot \mathbf{r}_{new} = 0$$




Figure 5.1: Interpretation of the orthogonality condition.

# Orthogonalization Schemes: Arnoldi and GMRES

## Arnoldi

- Uses Gram-Schmidt procedure to build an orthonormal basis

### ALGORITHM 6.4 Full Orthogonalization Method (FOM)

```

1.   Compute  $r_0 = b - Ax_0$ ,  $\beta := \|r_0\|_2$ , and  $v_1 := r_0/\beta$ 
2.   Define the  $m \times m$  matrix  $H_m = \{h_{ij}\}_{i,j=1,\dots,m}$ ; Set  $H_m = 0$ 
3.   For  $j = 1, 2, \dots, m$  Do:
4.       Compute  $w_j := Av_j$ 
5.       For  $i = 1, \dots, j$  Do:
6.            $h_{ij} = (w_j, v_i)$ 
7.            $w_j := w_j - h_{ij}v_i$ 
8.       EndDo
9.       Compute  $h_{j+1,j} = \|w_j\|_2$ . If  $h_{j+1,j} = 0$  set  $m := j$  and Goto 12
10.      Compute  $v_{j+1} = w_j/h_{j+1,j}$ .
11.  EndDo
12.  Compute  $y_m = H_m^{-1}(\beta e_1)$  and  $x_m = x_0 + V_m y_m$ 

```

## GMRES

- Generalized Minimum Residual

### ALGORITHM 6.9 GMRES

```

1.   Compute  $r_0 = b - Ax_0$ ,  $\beta := \|r_0\|_2$ , and  $v_1 := r_0/\beta$ 
2.   For  $j = 1, 2, \dots, m$  Do:
3.       Compute  $w_j := Av_j$ 
4.       For  $i = 1, \dots, j$  Do:
5.            $h_{ij} = (w_j, v_i)$ 
6.            $w_j := w_j - h_{ij}v_i$ 
7.       EndDo
8.        $h_{j+1,j} = \|w_j\|_2$ . If  $h_{j+1,j} = 0$  set  $m := j$  and go to 11
9.        $v_{j+1} = w_j/h_{j+1,j}$ 
10.  EndDo
11.  Define the  $(m+1) \times m$  Hessenberg matrix  $\bar{H}_m = \{h_{ij}\}_{1 \leq i \leq m+1, 1 \leq j \leq m}$ .
12.  Compute  $y_m$  the minimizer of  $\|\beta e_1 - \bar{H}_m y\|_2$  and  $x_m = x_0 + V_m y_m$ .

```



# Special cases: Symmetric Matrices

## Lanczos algorithm for symmetric matrices

- This is a simplification of Arnoldi's method, where the Hessenberg matrix is tridiagonal

### ALGORITHM 6.15 *The Lanczos Algorithm*

1. Choose an initial vector  $v_1$  of 2-norm unity. Set  $\beta_1 \equiv 0, v_0 \equiv 0$
2. For  $j = 1, 2, \dots, m$  Do:
3.      $w_j := Av_j - \beta_j v_{j-1}$  (If  $j = 1$  set  $\beta_1 v_0 \equiv 0$ )
4.      $\alpha_j := (w_j, v_j)$
5.      $w_j := w_j - \alpha_j v_j$
6.      $\beta_{j+1} := \|w_j\|_2$ . If  $\beta_{j+1} = 0$  then Stop
7.      $v_{j+1} := w_j / \beta_{j+1}$
8. EndDo

## Conjugate Gradient (CG)

- Only for symmetric systems

### ALGORITHM 6.16 *Lanczos Method for Linear Systems*

1. Compute  $r_0 = b - Ax_0$ ,  $\beta := \|r_0\|_2$ , and  $v_1 := r_0 / \beta$
2. For  $j = 1, 2, \dots, m$  Do:
3.      $w_j = Av_j - \beta_j v_{j-1}$  (If  $j = 1$  set  $\beta_1 v_0 \equiv 0$ )
4.      $\alpha_j = (w_j, v_j)$
5.      $w_j := w_j - \alpha_j v_j$
6.      $\beta_{j+1} = \|w_j\|_2$ . If  $\beta_{j+1} = 0$  set  $m := j$  and go to 9
7.      $v_{j+1} = w_j / \beta_{j+1}$
8. EndDo
9. Set  $T_m = \text{tridiag}(\beta_i, \alpha_i, \beta_{i+1})$ , and  $V_m = [v_1, \dots, v_m]$ .
10. Compute  $y_m = T_m^{-1}(\beta e_1)$  and  $x_m = x_0 + V_m y_m$

# Bi-Orthogonalization Schemes: Lanczos and BiCGSTAB

## Lanczos (Generalize for non-symmetric)

- Build 2 subspaces

$$\mathcal{K}_m(A, v_1) = \text{span}\{v_1, Av_1, \dots, A^{m-1}v_1\}$$

$$\mathcal{K}_m(A^T, w_1) = \text{span}\{w_1, A^T w_1, \dots, (A^T)^{m-1}w_1\}$$

### ALGORITHM 7.1 The Lanczos Biorthogonalization Procedure

1. Choose two vectors  $v_1, w_1$  such that  $(v_1, w_1) = 1$ .
2. Set  $\beta_1 = \delta_1 \equiv 0, w_0 = v_0 \equiv 0$
3. For  $j = 1, 2, \dots, m$  Do:
4.      $\alpha_j = (Av_j, w_j)$
5.      $\hat{v}_{j+1} = Av_j - \alpha_j v_j - \beta_j v_{j-1}$
6.      $\hat{w}_{j+1} = A^T w_j - \alpha_j w_j - \delta_j w_{j-1}$
7.      $\delta_{j+1} = |(\hat{v}_{j+1}, \hat{w}_{j+1})|^{1/2}$ . If  $\delta_{j+1} = 0$  Stop
8.      $\beta_{j+1} = (\hat{v}_{j+1}, \hat{w}_{j+1})/\delta_{j+1}$
9.      $w_{j+1} = \hat{w}_{j+1}/\beta_{j+1}$
10.     $v_{j+1} = \hat{v}_{j+1}/\delta_{j+1}$
11. EndDo

## BiCGSTAB

- Variation of Conjugate Gradient Squared (Transpose free Bi-Conjugate Gradient)

### ALGORITHM 7.7 BICGSTAB

1. Compute  $r_0 := b - Ax_0; r_0^*$  arbitrary;
2.  $p_0 := r_0$ .
3. For  $j = 0, 1, \dots$ , until convergence Do:
4.      $\alpha_j := (r_j, r_0^*)/(Ap_j, r_0^*)$
5.      $s_j := r_j - \alpha_j Ap_j$
6.      $\omega_j := (As_j, s_j)/(As_j, As_j)$
7.      $x_{j+1} := x_j + \alpha_j p_j + \omega_j s_j$
8.      $r_{j+1} := s_j - \omega_j As_j$
9.      $\beta_j := \frac{(r_{j+1}, r_0^*)}{(r_j, r_0^*)} \times \frac{\alpha_j}{\omega_j}$
10.     $p_{j+1} := r_{j+1} + \beta_j(p_j - \omega_j Ap_j)$
11. EndDo




# Convergence of Krylov Methods

- Convergence is related to properties of the eigenspectrum of the coefficient matrix,  $\mathbf{A}$ .

- But more specifically the condition number:

$$\|\mathbf{r}^{(\ell)}\| \propto \kappa(\mathbf{A})$$

- Convergence is often not monotonic
- GMRES does well ~~if~~ eigenvalues of  $\mathbf{A}$  are clustered.
- BiCGSTAB does well if eigenvalues are spread out
  - It isolates extremal eigenvalues





## More about condition numbers

- Generally it tells us how much an output value can change relate to a small change in the input
  - Bounds accuracy of approximate solution to a linear system
- Generally given by:

$$\kappa(\mathbf{A}) = \frac{\sigma_{\max}(\mathbf{A})}{\sigma_{\min}(\mathbf{A})}$$

- Recall that under certain conditions the SVD may be equivalent to the eigendecomposition so sometimes:

$$\kappa(\mathbf{A}) = \frac{|\lambda_{\max}(\mathbf{A})|}{|\lambda_{\min}(\mathbf{A})|}$$

- Note that this differs from the classical iterative techniques where generally convergence is given by:

$$\rho(\mathbf{A}) = \max_i |\lambda_i(\mathbf{A})| = |\lambda_{\max}(\mathbf{A})|$$

- The discretization of most PDE's give condition numbers that are ***unbounded***, and *increase as the problem size increases*.



# Preconditioning

- Objective: Lower the condition number of  $\mathbf{A}$  to achieve faster convergence
  - Left Preconditioner:  $\mathbf{M}^{-1}\mathbf{A}\mathbf{x} = \mathbf{M}^{-1}\mathbf{b}$
  - Right Preconditioner:  $\mathbf{A}\mathbf{M}^{-1}\mathbf{M}\mathbf{x} = \mathbf{b}$  solve  $\mathbf{A}\mathbf{M}^{-1}\mathbf{u} = \mathbf{b}$ , then  $\mathbf{u} = \mathbf{M}\mathbf{x}$
- Properties of a good preconditioner
  - Large decrease in condition number for a wide range of  $\mathbf{A}$
  - Efficient to construct and apply
- Krylov methods generally only as good as their preconditioner.
  - This is conservation of misery.
  - Research into preconditioners still an active topic
  - Consequently, for some problems classical iteration schemes are still the best choice



# Lots of Types of Preconditioners

- Classical methods; Jacobi, SOR, SSOR
- Decomposition:
  - Incomplete LU (ILU): do an LU factorization but don't change sparsity
    - Very difficult to parallelize
  - Cholesky
- Polynomials: Chebyshev, Least-Squares
- Multigrid
- “Physics-based”



# Summary of Krylov Methods

## Orthogonalization

- Usually requires more storage
  - e.g. Hessenberg matrix and GMRES vectors
  - Restarted forms attempt to mitigate this
- Guaranteed to converge
  - May require full orthogonalization

## Bi-Orthogonalization

- Avoid storage by doing twice the work (bi-orthogonalization)
  - Uses 3-term recurrence for orthogonalizing
- Not guaranteed to converge, but it usually does

**Probably need a preconditioner**



# Scientific Computing Libraries

BLAS, LAPACK, PETSc, Trilinos, and Others



# BLAS (Basic Linear Algebra Subprograms)

- BLAS is a (Fortran) programming interface to low-level linear algebra routines.
- Examples of basic linear algebra operations:
  - Dot products, addition of vectors, scalar multiplication of vectors.
  - Matrix-vector multiplications.
  - Matrix-matrix multiplications.
- Why BLAS?
  - Handwritten simple linear algebra routines can run at widely varying speeds.
  - Loop unrolling and finding correct compiler flags is key and sometimes becomes difficult when using a new compiler or computer.
  - Basic linear algebra routines form the backbone of many sophisticated solvers and BLAS package tries to provide most optimized version of basic linear algebra routines.
  - Linking to BLAS, we get code that runs much faster than hand-written version of code.





# BLAS (Basic Linear Algebra Subprograms)

- BLAS-1 operations:
  - Routines which involve only vector operations: dot-products, vector norms.
  - S – Single Precision, D - Double Precision, C – Complex, Z – Double precision complex

<code>_SWAP</code>	<code>x &lt;---&gt; y</code>	<code>S, D</code>
<code>_SCAL</code>	<code>x &lt;-- alpha * x</code>	<code>S, D, C, Z, CS, ZD</code>
<code>_COPY</code>	<code>x &lt;-- y</code>	<code>S, D, C, Z</code>
<code>_AXPY</code>	<code>y &lt;-- alpha * x + y</code>	<code>S, D, C, Z</code>
<code>_DOT</code>	<code>dot &lt;--</code>	<code>S, D, DS</code>
<code>_DOTU</code>	<code>dot &lt;-- x^T*y</code>	<code>C, Z</code>
<code>_DOTC</code>	<code>dot &lt;-- x^H*y</code>	<code>C, Z</code>
<code>_NRM2</code>	<code>nrm2 &lt;--    x   _2</code>	<code>S, D, SC, DZ</code>
<code>_ASUM</code>	<code>nrm1 &lt;--    x   _1</code>	<code>S, D, SC, DZ</code>





# BLAS (Basic Linear Algebra Subprograms)

- BLAS-2 operations:
  - This level contains matrix-vector operations including, among other things, a generalized matrix-vector multiplication (gemv):

<code>_GEMV</code>	<code>y &lt;-- alpha*A*x + beta*y, y &lt;-- alpha*A^T*x + beta*y</code>	(General Real Matrix)
<code>_GBMV</code>	<code>y &lt;-- alpha*A*x + beta*y, y &lt;-- alpha*A^T*x + beta*y</code>	(General Banded)
<code>_HEMV</code>	<code>y &lt;-- alpha*A*x + beta*y</code>	(Hermitian (complex))
<code>_HBMV</code>	<code>y &lt;-- alpha*A*x + beta*y</code>	(Hermitian (banded))
<code>_SYMV</code>	<code>y &lt;-- alpha*A*x + beta*y</code>	(Symmetric)
<code>_SBMV</code>	<code>y &lt;-- alpha*A*x + beta*y</code>	(Symmetric banded)
<code>_TRMV</code>	<code>y &lt;-- A*x , x &lt;-- A^T*x</code>	(Triangular)
<code>_TBMV</code>	<code>y &lt;-- A*x , x &lt;-- A^T*x</code>	(Triangular banded)
<code>_TRSV</code>	<code>y &lt;-- inv(A)*x , x &lt;-- inv(A^T) * x</code>	



# BLAS (Basic Linear Algebra Subprograms)

- BLAS-3 Matrix-matrix operations:
  - This level operations are matrix-matrix multiplications.

```
_GEMM  C <-- alpha op(A) op(B) + beta C  
_SYMM  C <-- alpha AB + beta C  
_HEMM  C <-- alpha AB + beta C  
_SYRK  C <-- alpha A A^T + beta C  
_HERK  C <-- alpha A A^H + beta C  
_SYRK2 C <-- alpha A B^T + alpha B A^T + beta C  
_TRMM  B <-- alpha op(A) B  
_TRSM  B <-- alpha op(inv(A)) B
```



# BLAS-Example Code

```
51 * =====
52 * DOUBLE PRECISION FUNCTION ddot(N,DX,INCX,DY,INCY)
53 *
54 * -- Reference BLAS level1 routine (version 3.4.0) --
55 * -- Reference BLAS is a software package provided by Univ. of Tennessee, --
56 * -- Univ. of California Berkeley, Univ. of Colorado Denver and NAG Ltd.--
57 * November 2011
58 *
59 * .. Scalar Arguments ..
60 * INTEGER INCX,INCY,N
61 * ..
62 * .. Array Arguments ..
63 * DOUBLE PRECISION DX(*),DY(*)
64 * ..
65 *
66 * =====
67 *
68 * .. Local Scalars ..
69 * DOUBLE PRECISION DTEMP
70 * INTEGER I,IX,IY,M,MP1
71 * ..
72 * .. Intrinsic Functions ..
73 * INTRINSIC mod
74 * ..
75 * ddot = 0.0d0
76 * dtemp = 0.0d0
77 * IF (n.LE.0) RETURN
78 * IF (incx.EQ.1 .AND. incy.EQ.1) THEN
79 *
80 *     code for both increments equal to 1
81 *
82 *
83 *     clean-up loop
84 *
85 *     m = mod(n,5)
86 *     IF (m.NE.0) THEN
87 *         DO i = 1,m
88 *             dtemp = dtemp + dx(i)*dy(i)
89 *         END DO
90 *         IF (n.LT.5) THEN
91 *             ddot=dtemp
92 *             RETURN
93 *         END IF
94 *     END IF
95 *     mp1 = m + 1
96 *     DO i = mp1,n,5
97 *         dtemp = dtemp + dx(i)*dy(i) + dx(i+1)*dy(i+1) +
98 *             $ dx(i+2)*dy(i+2) + dx(i+3)*dy(i+3) + dx(i+4)*dy(i+4)
99 *     END DO
100 * ELSE
101 * 
```



# LAPACK

- LA Pack is a collection of Fortran functions that can help you solve Linear Algebra related problems based on BLAS routines.
- LAPACK provides routines for solving systems of simultaneous linear equations, least-squares solutions of linear systems of equations, eigenvalue problems, and singular value problems.
- The associated matrix factorizations (LU, Cholesky, QR, SVD, Schur, generalized Schur) are also provided,
- Dense and banded matrices are handled, but not general sparse matrices. In all areas, similar functionality is provided for real and complex matrices, in both single and double precision.
- Online Reference: <http://www.netlib.org/lapack>



## LAPACK Routines

- Simple driver routines: Simple driver routines solve a linear algebra problem
- Eg: Finding eigenvalues of a matrix, solving a set of linear equations etc.,
- Expert driver routines: Expert driver routines do the same things as simple driver routines, but will provide more options or information to the user.
- Eg: SGESV is used to solve linear systems whereas the expert driver SGESVX not only solves the linear system but will also provide the estimate of the condition number of input matrix.
- Computational routines: Routines are mainly for internal use by LA Pack itself and called by driver routines.
- Eg: LU, QR and other factorizations or reduction of symmetric matrix to tridiagonal form.



# LAPACK Naming conventions

- Lapack functions are usually named in the form **XYZZZ**

**X = type of problem that the routine solves**

**S Single precision real**

**D Double precision real**

**C Single complex**

**Z Double precision complex**

**YY = matrix types**

**GE General**

**BD Bidiagonal**

**HE Hermitian**

**HB Hermitian Band**

**SB Symmetric Band**

**ZZZ = indicate the computation performed**

**Eg: SV = Solve, SVX = Solve Expert**



# PETSc

- Portable Extensible Toolkit for Scientific computing:

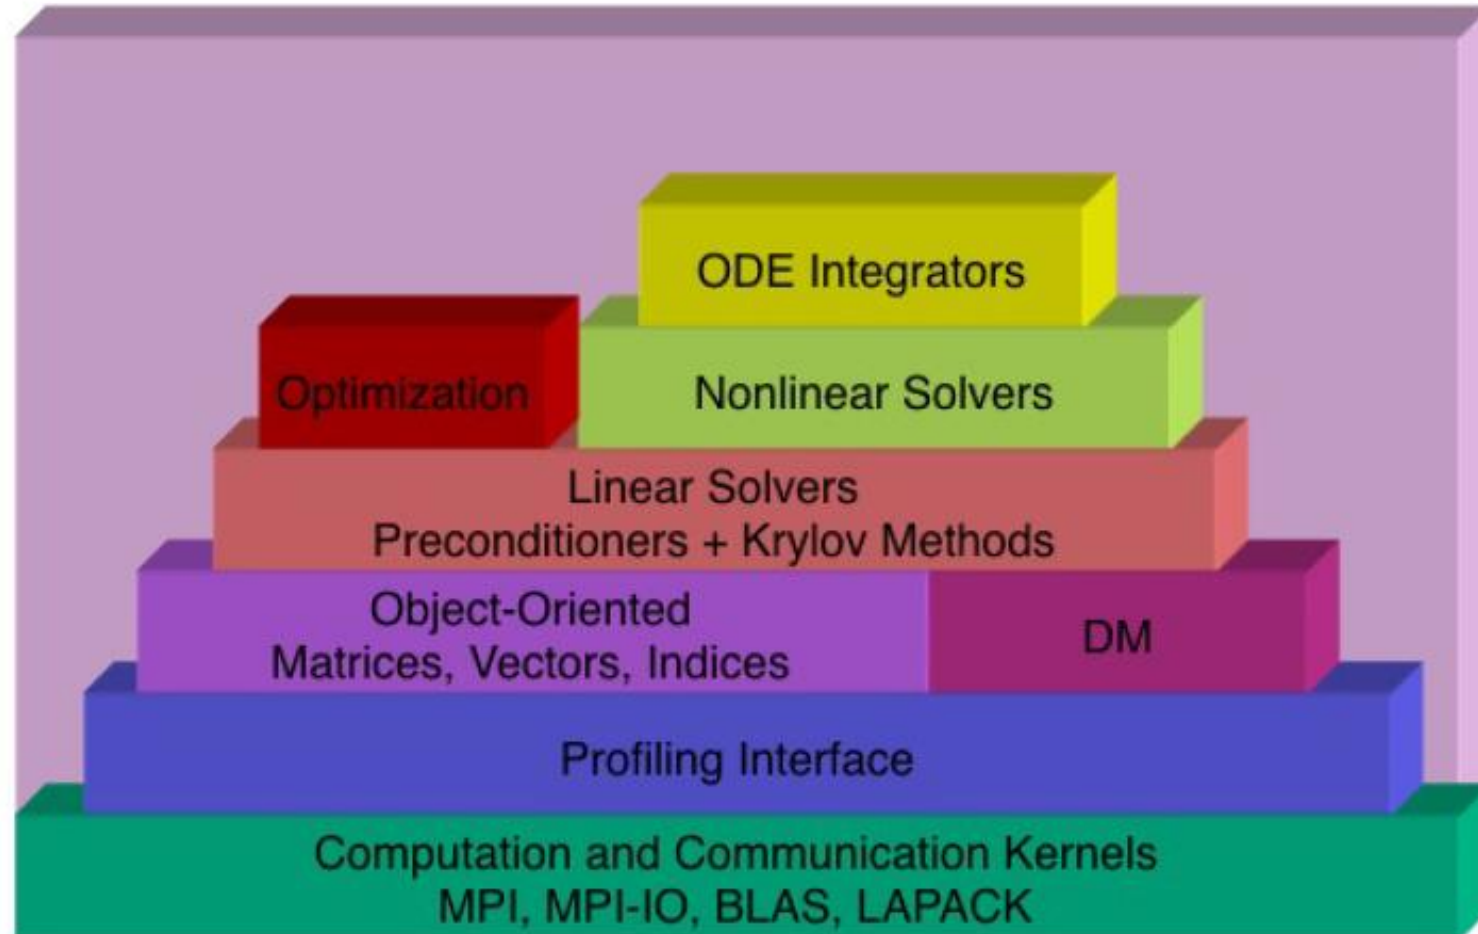
*Developing parallel, nontrivial PDE solvers that deliver high performance is still difficult and requires months (or even years) of concentrated effort. PETSc is a toolkit that can ease these difficulties and reduce the development time, but it is not a black-box PDE solver, nor a silver bullet.* — Barry Smith

- Philosophy: Everything has a plugin architecture
  - Vectors, Matrices, Partitioning algorithms
  - Preconditioners, Krylov accelerators
  - Nonlinear solvers, Time integrators
  - Spatial discretizations
  - Application user loads plugin at run, no source code in sight. PETSc tries to keep solvers independent of physics and discretization.





## PETSc Structure





# Basic PETSc object usage

Every object in PETSc supports a basic interface

Function	Operation
<code>Create()</code>	create the object
<code>Get/SetName()</code>	name the object
<code>Get/SetType()</code>	set the implementation type
<code>Get/SetOptionsPrefix()</code>	set the prefix for all options
<code>SetFromOptions()</code>	customize object from the command line
<code>SetUp()</code>	perform other initialization
<code>View()</code>	view the object
<code>Destroy()</code>	cleanup object allocation

Also, all objects support the `-help` option.



## PETSc Vectors

- PETSc vectors are fundamental datatypes of PETSc which is used represent field solutions, right-hand sides etc. Each process locally owns a subvector of contiguous global data.
- Creating PETSc vectors:
  - `VecCreate(MPI_Comm, Vec *)`
  - `VecSetSizes(Vec, int n, int N)`
  - `VecSetTypes(Vec, VecType typename)`
  - `VecSetFromOptions(Vec)`
- Supports all vector space operations `VecDot()`, `VecNorm()`, `VecScale()`



## PETSc Vectors

- Inserting entries into PETSc vectors:
  - Each process sets or add values and begins communications to send values to correct process and complete the communication.
  - `VecSetValues(Vec v, int n, int rows[], PetscScalar values[], mode)`
  - `VecAssemblyBegin(Vec v)`
  - `VecAssemblyEnd(Vec v)`
- PETSc allows you to access the local storage with `VecGetArray()` functions.


# PETSc Matrices

- PETSc handles both sparse and dense matrix formats in parallel.

```
MatCreate (MPI_Comm, Mat *)  
MatSetSizes (Mat, int m, int n, int M, int N)  
MatSetType (Mat, MatType typeName)  
MatSetFromOptions (Mat)  
    Can set the type at runtime  
MatMPIBAIJSetPreallocation (Mat, ...)  
    important for assembly performance  
MatSetBlockSize (Mat, int bs)  
    for vector problems  
MatSetValues (Mat, ...)  
    MUST be used, but does automatic communication  
    MatSetValuesLocal, MatSetValuesStencil,  
    MatSetValuesBlocked
```

# Parallel sparse matrix in PETSc

Each process locally owns a submatrix of contiguous global rows  
 Each submatrix consists of diagonal and off-diagonal parts



`MatGetOwnershipRange (Mat A, int *start, int *end)`  
 start: first locally owned row of global matrix  
 end-1: last locally owned row of global matrix



## Other useful features of PETSc

- Iterative solvers:
  - Linear solvers in PETSc KSP: Conjugate Gradient, Bi Conjugate Gradient, GMRES, etc.
  - Lots of sophisticated Preconditioners like block Jacobi, SOR, Multigrid, field-split, etc.
  - Nonlinear solvers (SNES):
    - Newton type with line search and trust-region
    - Quasi Newton methods
    - Nonlinear conjugate gradients
    - User-defined methods.
- Time Integration strategies



# TRILINOS

- The Trilinos Project is an effort to develop algorithms and enabling technologies within an object-oriented software framework for the solution of large-scale, complex multi-physics engineering and scientific problems. A unique design feature of Trilinos is its focus on packages.
- More extensive than PETSc and more complex to use.
- Trilinos tries to provide an environment for solving FEM problems and PETSc provides an environment for solving sparse linear algebra problems.

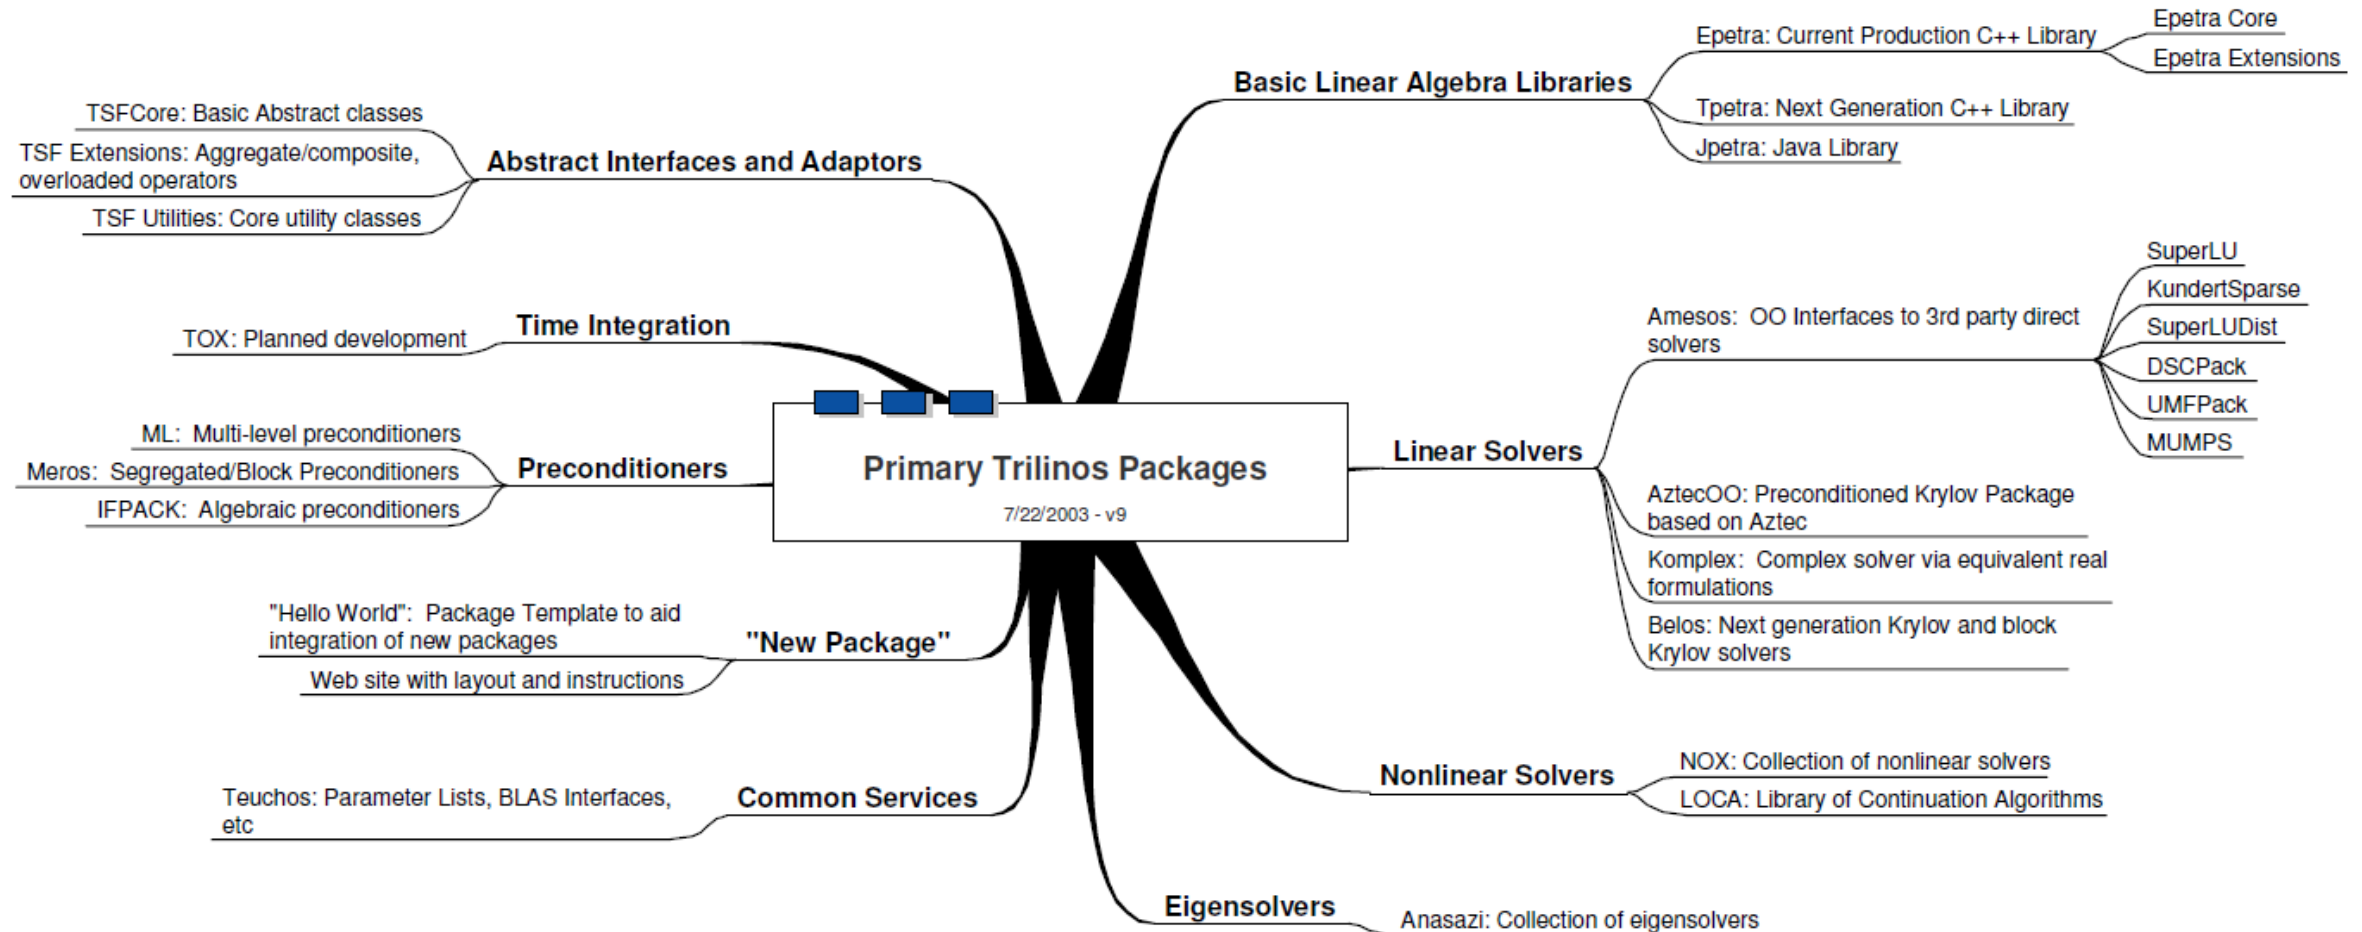


# Trilinos Capabilities

- Trilinos – Greek for “String of Pearls”
  - Most package names are Greek
  - Duplication of capability
  - Some are deprecated



User Experience	Parallel Programming Environments	I/O Support
Mesh & Geometry	Framework & Tools	Discretization
Linear Algebra Services	Linear & Eigen Solvers	Embedded Nonlinear Analysis
Software Engineering		





- Contains Tools for:
  - Problem Discretization
  - Solution of Algebraic Systems
  - Uncertainty Quantification
  - Numerical Optimization

<https://fastmath-scidac.llnl.gov/software-catalog.html>

# Integrating Libraries with your Code

## xSDK Version 0.4.0: December 2018

